

Mercedes-Benz Greener Manufacturing

August 19, 2023

```
[1]: import pandas as pd
import math
import matplotlib.pyplot as plt
import seaborn as sns
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn import preprocessing
from sklearn.preprocessing import LabelEncoder
%matplotlib inline
```

```
[2]: df_test = pd.read_csv('test.csv')
df_test.head()
```

```
[2]:   ID  X0 X1  X2 X3 X4 X5 X6 X8  X10  ...  X375  X376  X377  X378  X379  X380  \
0    1  az  v   n  f  d  t  a  w    0  ...    0     0     0     1     0     0
1    2   t  b  ai  a  d  b  g  y    0  ...    0     0     1     0     0     0
2    3  az  v  as  f  d  a  j  j    0  ...    0     0     0     1     0     0
3    4  az  l   n  f  d  z  l  n    0  ...    0     0     0     1     0     0
4    5   w  s  as  c  d  y  i  m    0  ...    1     0     0     0     0     0
```

```
      X382  X383  X384  X385
0         0     0     0     0
1         0     0     0     0
2         0     0     0     0
3         0     0     0     0
4         0     0     0     0
```

[5 rows x 377 columns]

```
[3]: df_train = pd.read_csv("train.csv")
df_train.head()
```

```
[3]:   ID      y  X0 X1  X2 X3 X4 X5 X6 X8  ...  X375  X376  X377  X378  X379  \
0    0  130.81  k  v  at  a  d  u  j  o  ...    0     0     1     0     0
1    6   88.53  k  t  av  e  d  y  l  o  ...    1     0     0     0     0
2    7   76.26  az  w   n  c  d  x  j  x  ...    0     0     0     0     0
3    9   80.62  az  t   n  f  d  x  l  e  ...    0     0     0     0     0
```

4	13	78.02	az	v	n	f	d	h	d	n	...	0	0	0	0	0
---	----	-------	----	---	---	---	---	---	---	---	-----	---	---	---	---	---

	X380	X382	X383	X384	X385
0	0	0	0	0	0
1	0	0	0	0	0
2	0	1	0	0	0
3	0	0	0	0	0
4	0	0	0	0	0

[5 rows x 378 columns]

```
[4]: ## lets print the shape,column nd info for the train and test data of
      ↳mercedes_benz
print("the shape of the train set",df_train.shape)
print("the shape of test set",df_test.shape)
```

```
the shape of the train set (4209, 378)
the shape of test set (4209, 377)
```

```
[5]: df_train.columns
```

```
[5]: Index(['ID', 'y', 'X0', 'X1', 'X2', 'X3', 'X4', 'X5', 'X6', 'X8',
          ...
          'X375', 'X376', 'X377', 'X378', 'X379', 'X380', 'X382', 'X383', 'X384',
          'X385'],
          dtype='object', length=378)
```

```
[6]: df_test.columns
```

```
[6]: Index(['ID', 'X0', 'X1', 'X2', 'X3', 'X4', 'X5', 'X6', 'X8', 'X10',
          ...
          'X375', 'X376', 'X377', 'X378', 'X379', 'X380', 'X382', 'X383', 'X384',
          'X385'],
          dtype='object', length=377)
```

```
[7]: df_train.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4209 entries, 0 to 4208
Columns: 378 entries, ID to X385
dtypes: float64(1), int64(369), object(8)
memory usage: 12.1+ MB
```

```
[8]: df_test.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4209 entries, 0 to 4208
Columns: 377 entries, ID to X385
```

```
dtypes: int64(369), object(8)
memory usage: 12.1+ MB
```

```
[9]: df_train.describe()
```

```
[9]:
```

	ID	y	X10	X11	X12	\
count	4209.000000	4209.000000	4209.000000	4209.0	4209.000000	
mean	4205.960798	100.669318	0.013305	0.0	0.075077	
std	2437.608688	12.679381	0.114590	0.0	0.263547	
min	0.000000	72.110000	0.000000	0.0	0.000000	
25%	2095.000000	90.820000	0.000000	0.0	0.000000	
50%	4220.000000	99.150000	0.000000	0.0	0.000000	
75%	6314.000000	109.010000	0.000000	0.0	0.000000	
max	8417.000000	265.320000	1.000000	0.0	1.000000	

	X13	X14	X15	X16	X17	...	\
count	4209.000000	4209.000000	4209.000000	4209.000000	4209.000000	...	
mean	0.057971	0.428130	0.000475	0.002613	0.007603	...	
std	0.233716	0.494867	0.021796	0.051061	0.086872	...	
min	0.000000	0.000000	0.000000	0.000000	0.000000	...	
25%	0.000000	0.000000	0.000000	0.000000	0.000000	...	
50%	0.000000	0.000000	0.000000	0.000000	0.000000	...	
75%	0.000000	1.000000	0.000000	0.000000	0.000000	...	
max	1.000000	1.000000	1.000000	1.000000	1.000000	...	

	X375	X376	X377	X378	X379	\
count	4209.000000	4209.000000	4209.000000	4209.000000	4209.000000	
mean	0.318841	0.057258	0.314802	0.020670	0.009503	
std	0.466082	0.232363	0.464492	0.142294	0.097033	
min	0.000000	0.000000	0.000000	0.000000	0.000000	
25%	0.000000	0.000000	0.000000	0.000000	0.000000	
50%	0.000000	0.000000	0.000000	0.000000	0.000000	
75%	1.000000	0.000000	1.000000	0.000000	0.000000	
max	1.000000	1.000000	1.000000	1.000000	1.000000	

	X380	X382	X383	X384	X385
count	4209.000000	4209.000000	4209.000000	4209.000000	4209.000000
mean	0.008078	0.007603	0.001663	0.000475	0.001426
std	0.089524	0.086872	0.040752	0.021796	0.037734
min	0.000000	0.000000	0.000000	0.000000	0.000000
25%	0.000000	0.000000	0.000000	0.000000	0.000000
50%	0.000000	0.000000	0.000000	0.000000	0.000000
75%	0.000000	0.000000	0.000000	0.000000	0.000000
max	1.000000	1.000000	1.000000	1.000000	1.000000

```
[8 rows x 370 columns]
```

```
[10]: df_test.describe()
```

```
[10]:
```

	ID	X10	X11	X12	X13	\
count	4209.000000	4209.000000	4209.000000	4209.000000	4209.000000	
mean	4211.039202	0.019007	0.000238	0.074364	0.061060	
std	2423.078926	0.136565	0.015414	0.262394	0.239468	
min	1.000000	0.000000	0.000000	0.000000	0.000000	
25%	2115.000000	0.000000	0.000000	0.000000	0.000000	
50%	4202.000000	0.000000	0.000000	0.000000	0.000000	
75%	6310.000000	0.000000	0.000000	0.000000	0.000000	
max	8416.000000	1.000000	1.000000	1.000000	1.000000	

	X14	X15	X16	X17	X18	...	\
count	4209.000000	4209.000000	4209.000000	4209.000000	4209.000000	...	
mean	0.427893	0.000713	0.002613	0.008791	0.010216	...	
std	0.494832	0.026691	0.051061	0.093357	0.100570	...	
min	0.000000	0.000000	0.000000	0.000000	0.000000	...	
25%	0.000000	0.000000	0.000000	0.000000	0.000000	...	
50%	0.000000	0.000000	0.000000	0.000000	0.000000	...	
75%	1.000000	0.000000	0.000000	0.000000	0.000000	...	
max	1.000000	1.000000	1.000000	1.000000	1.000000	...	

	X375	X376	X377	X378	X379	\
count	4209.000000	4209.000000	4209.000000	4209.000000	4209.000000	
mean	0.325968	0.049656	0.311951	0.019244	0.011879	
std	0.468791	0.217258	0.463345	0.137399	0.108356	
min	0.000000	0.000000	0.000000	0.000000	0.000000	
25%	0.000000	0.000000	0.000000	0.000000	0.000000	
50%	0.000000	0.000000	0.000000	0.000000	0.000000	
75%	1.000000	0.000000	1.000000	0.000000	0.000000	
max	1.000000	1.000000	1.000000	1.000000	1.000000	

	X380	X382	X383	X384	X385
count	4209.000000	4209.000000	4209.000000	4209.000000	4209.000000
mean	0.008078	0.008791	0.000475	0.000713	0.001663
std	0.089524	0.093357	0.021796	0.026691	0.040752
min	0.000000	0.000000	0.000000	0.000000	0.000000
25%	0.000000	0.000000	0.000000	0.000000	0.000000
50%	0.000000	0.000000	0.000000	0.000000	0.000000
75%	0.000000	0.000000	0.000000	0.000000	0.000000
max	1.000000	1.000000	1.000000	1.000000	1.000000

```
[8 rows x 369 columns]
```

0.0.1 If for any column(s) the variance is equal to zero then you need to remove those variable.

```
[11]: df_train.var()
```

```
[11]: ID      5.941936e+06
      y      1.607667e+02
      X10    1.313092e-02
      X11    0.000000e+00
      X12    6.945713e-02
      ...
      X380    8.014579e-03
      X382    7.546747e-03
      X383    1.660732e-03
      X384    4.750593e-04
      X385    1.423823e-03
      Length: 370, dtype: float64
```

```
[12]: ## check if the train set variance is equal t zero
      df_train.var()==0
```

```
[12]: ID      False
      y      False
      X10    False
      X11     True
      X12    False
      ...
      X380    False
      X382    False
      X383    False
      X384    False
      X385    False
      Length: 370, dtype: bool
```

```
[13]: (df_train.var()==0).values
```

```
[13]: array([False, False, False,  True, False, False, False, False, False,
        False, False, False, False, False, False, False, False, False,
        False, False, False, False, False, False, False, False, False,
        False, False, False, False, False, False, False, False, False,
        False, False, False, False, False, False, False, False, False,
        False, False, False, False, False, False, False, False, False,
        False, False, False, False, False, False, False, False, False,
        False, False,  True, False, False, False, False, False, False,
        False, False, False, False, False, False, False, False,  True, False,
```

```
False, False, False, False, False, False, False, False, False,
False, False, False, False, False, False, False, False, False,
False, False, False, False, False, False, False, False, False,
False, False, False, False, False, False, False, False, False,
False, False, False, False, False, False, False, False, False,
False, False, False, False, False, False, False, False, False,
False, False, False, False, False, False, False, False, False,
False, False, False, False, False, False, False, False, False,
False, False, False, False, False, False, False, False, False,
False, False, False, False, False, False, False, False, False,
False, False, False, False, False, False, False, False, False,
False, False, False, True, False, True, False, False, False,
False, False, False, False, False, False, False, False, False,
False, False, False, False, False, False, False, False, False,
False, False, False, False, False, False, False, False, False,
False, False, True, False, False, False, False, False, False,
False, False, False, False, False, False, False, False, False,
False, False, False, False, False, True, True, False, False,
True, False, False, False, True, False, False, False, False,
False, False, False, False, False, False, False, False, False,
False, False, False, False, False, False, False, False, False,
False, False, False, False, False, False, False, False, False,
True, False, False, False, False, False, False, False, False,
False, False, False, False, False, False, False, False, True,
False, False, False, False, False, False, False, False, False,
False, False, False, False, False, False, False, False, False,
False, False, False, False, False, False, False, False, False,
False, False, False, False, False, False, False, False, False,
False])
```

```
[14]: train_variance_zero =df_train.var()[df_train.var()==0].index.values
train_variance_zero
```

```
[14]: array(['X11', 'X93', 'X107', 'X233', 'X235', 'X268', 'X289', 'X290',
'X293', 'X297', 'X330', 'X347'], dtype=object)
```

```
[15]: ## removing the data whose variance is equal to zero
print("before dropping the vriance with zero the train set",df_train.shape)
df_train= df_train.drop(train_variance_zero,axis=1)
print("Eliminating the zero variance columns",df_train.shape)
```

before dropping the vriance with zero the train set (4209, 378)
Eliminating the zero variance columns (4209, 366)

In the bove we can see that the 12 columns has been dropped from the original columns whose variance is zero.

```
## Similarly check for test data
df_test.var()
```

```
ID          5.871311e+06
X10         1.865006e-02
X11         2.375861e-04
X12         6.885074e-02
X13         5.734498e-02
...
X380        8.014579e-03
X382        8.715481e-03
X383        4.750593e-04
X384        7.124196e-04
X385        1.660732e-03
Length: 369, dtype: float64
```

```
(df_test.var()==0).values
```

[illegible]

```
False, False, False, False, False, False, False, False, False,
False, False, False, False, False, False, False, False, False,
False, True, True, False, False, False, False, False, False,
False, False, False, False, False, False, False, False, False,
False, False, False, False, False, False, False, False, False,
False, False, False, False, False, False, False, False, False,
False, False, False, False, False, False, False, False, False,
False, False, False, False, False, False, False, False, False,
False, False, False, False, False, False, False, False, False,
False, False, False, False, False, False, False, False, False,
False, False, True, False, False, False, False, False, False,
False, False, False, False, False, False, False, False, False]
```

```
[18]: test_variance_zero = df_test.var()[df_test.var()==0].index.values
test_variance_zero
```

```
[18]: array(['X257', 'X258', 'X295', 'X296', 'X369'], dtype=object)
```

```
[19]: print("The test data set with zero variance",df_test.shape)
df_test= df_test.drop(test_variance_zero,axis=1)
print("The test data without zero variance",df_test.shape)
```

The test data set with zero variance (4209, 377)

The test data without zero variance (4209, 372)

we can see that 5 columns has been eliminated whose variance is zero.

0.0.2 Check for null and unique values for test and train sets.

```
[20]: # lets check the relevant columns for our prediction
df_train.columns
```

```
[20]: Index(['ID', 'y', 'X0', 'X1', 'X2', 'X3', 'X4', 'X5', 'X6', 'X8',
...
'X375', 'X376', 'X377', 'X378', 'X379', 'X380', 'X382', 'X383', 'X384',
'X385'],
dtype='object', length=366)
```

```
[21]: df_test.columns
```

```
[21]: Index(['ID', 'X0', 'X1', 'X2', 'X3', 'X4', 'X5', 'X6', 'X8', 'X10',
...
'X375', 'X376', 'X377', 'X378', 'X379', 'X380', 'X382', 'X383', 'X384',
'X385'],
dtype='object', length=372)
```


we can see that ID columns has no relevancy in prediction of the data for sales,safety ,optimizing the speed.

```
[22]: df_train=df_train.drop(["ID"],axis=1)
df_train.head()
```

```
[22]:
```

	y	X0	X1	X2	X3	X4	X5	X6	X8	X10	...	X375	X376	X377	X378	X379	\
0	130.81	k	v	at	a	d	u	j	o	0	...	0	0	1	0	0	
1	88.53	k	t	av	e	d	y	l	o	0	...	1	0	0	0	0	
2	76.26	az	w	n	c	d	x	j	x	0	...	0	0	0	0	0	
3	80.62	az	t	n	f	d	x	l	e	0	...	0	0	0	0	0	
4	78.02	az	v	n	f	d	h	d	n	0	...	0	0	0	0	0	

	X380	X382	X383	X384	X385
0	0	0	0	0	0
1	0	0	0	0	0
2	0	1	0	0	0
3	0	0	0	0	0
4	0	0	0	0	0

[5 rows x 365 columns]

```
[23]: df_test=df_test.drop(["ID"],axis=1)
df_test.head()
```

```
[23]:
```

	X0	X1	X2	X3	X4	X5	X6	X8	X10	X11	...	X375	X376	X377	X378	X379	\
0	az	v	n	f	d	t	a	w	0	0	...	0	0	0	1	0	
1	t	b	ai	a	d	b	g	y	0	0	...	0	0	1	0	0	
2	az	v	as	f	d	a	j	j	0	0	...	0	0	0	1	0	
3	az	l	n	f	d	z	l	n	0	0	...	0	0	0	1	0	
4	w	s	as	c	d	y	i	m	0	0	...	1	0	0	0	0	

	X380	X382	X383	X384	X385
0	0	0	0	0	0
1	0	0	0	0	0
2	0	0	0	0	0
3	0	0	0	0	0
4	0	0	0	0	0

[5 rows x 371 columns]

```
[24]: # checking for the null values in train set and test set for mercedes_Benz
df_train.isnull().sum()
```

```
[24]: y      0
X0      0
X1      0
X2      0
```

X3	0
	..
X380	0
X382	0
X383	0
X384	0
X385	0

```
[25]: df_test.isnull().sum()
```

[25]:	X0	0
	X1	0
	X2	0
	X3	0
	X4	0
		..
	X380	0
	X382	0
	X383	0
	X384	0
	X385	0

```
[26]: df_train.isna().sum().values
```

[illegible]

```
[27]: df_test.isna().sum().values
```

```
[28]: # find the uniques values in test and train set
df_train.nunique()
```

```
[29]: df_train.nunique().values
```

11

[illegible]

0.0.3 Apply label encoder.

Before applying label encoder lets bifurcate the categorical variable (object dtypes).

```
[32]: cat_train = df_train.select_dtypes(include=[object])
      cat_train.head()
```

```
[32]:      X0 X1  X2 X3 X4 X5 X6 X8
      0   k  v   at  a  d  u   j  o
      1   k  t   av  e  d  y   l  o
      2  az  w    n  c  d  x   j  x
      3  az  t    n  f  d  x   l  e
      4  az  v    n  f  d  h   d  n
```

```
[33]: cat_test = df_test.select_dtypes(include=[object])
      cat_test.head()
```

```
[33]:      X0 X1  X2 X3 X4 X5 X6 X8
      0  az  v   n  f  d  t  a  w
      1   t  b  ai  a  d  b  g  y
      2  az  v  as  f  d  a  j  j
      3  az  l   n  f  d  z  l  n
      4   w  s  as  c  d  y  i  m
```

```
[34]: ## lets check the columns, describe and shape of the new train categorical data
cat train.columns
```

```
[34]: Index(['X0', 'X1', 'X2', 'X3', 'X4', 'X5', 'X6', 'X8'], dtype='object')
```

```
[35]: cat_train.shape
```

```
[35]: (4209, 8)
```

```
[36]: cat_train.describe()
```

```
[36]:
```

	X0	X1	X2	X3	X4	X5	X6	X8
count	4209	4209	4209	4209	4209	4209	4209	4209
unique	47	27	44	7	4	29	12	25
top	z	aa	as	c	d	v	g	j
freq	360	833	1659	1942	4205	231	1042	277

```
[37]: cat_test.columns
```

```
[37]: Index(['X0', 'X1', 'X2', 'X3', 'X4', 'X5', 'X6', 'X8'], dtype='object')
```

```
[38]: cat_test.shape
```

```
[38]: (4209, 8)
```

```
[39]: cat_test.describe()
```

```
[39]:
```

	X0	X1	X2	X3	X4	X5	X6	X8
count	4209	4209	4209	4209	4209	4209	4209	4209
unique	49	27	45	7	4	32	12	25
top	ak	aa	as	c	d	v	g	e
freq	432	826	1658	1900	4203	246	1073	274

```
[40]: le =LabelEncoder()
```

```
[41]: df_train['X0'] =le.fit_transform(df_train['X0'])
```

```
[42]: df_train.X0
```

```
[42]: 0      32
1      32
2      20
3      20
4      20
..
4204    8
4205   31
4206    8
4207    9
4208   46
Name: X0, Length: 4209, dtype: int64
```

```
[43]: # similarly we encode the data for columns X1,X2,X3,X4,X5,X6,X8
```

```
df_train['X1']= le.fit_transform(df_train['X1'])
df_train['X2']=le.fit_transform(df_train['X2'])
df_train['X3']=le.fit_transform(df_train['X3'])
df_train['X4']=le.fit_transform(df_train['X4'])
df_train['X5']=le.fit_transform(df_train['X5'])
df_train['X6']=le.fit_transform(df_train['X6'])
df_train['X8']=le.fit_transform(df_train['X8'])
```

```
[44]: df_test.X0
```

```
[44]: 0      az
      1      t
      2      az
      3      az
      4      w
      ..
     4204    aj
     4205      t
     4206      y
     4207    ak
     4208      t
      Name: X0, Length: 4209, dtype: object
```

```
[45]: # Similarly we can use label Encoders for test datasets
```

```
df_test['X0']= le.fit_transform(df_test['X0'])
df_test['X1']=le.fit_transform(df_test['X1'])
df_test['X2']=le.fit_transform(df_train['X2'])
df_test['X3']=le.fit_transform(df_test['X3'])
df_test['X4']=le.fit_transform(df_test['X4'])
df_test['X5']=le.fit_transform(df_train['X5'])
df_test['X6']=le.fit_transform(df_test['X6'])
df_test['X8']=le.fit_transform(df_test['X8'])
```

```
[46]: df_train.head()
```

```
[46]:
```

	y	X0	X1	X2	X3	X4	X5	X6	X8	X10	...	X375	X376	X377	X378	\
0	130.81	32	23	17	0	3	24	9	14	0	...	0	0	1	0	
1	88.53	32	21	19	4	3	28	11	14	0	...	1	0	0	0	
2	76.26	20	24	34	2	3	27	9	23	0	...	0	0	0	0	
3	80.62	20	21	34	5	3	27	11	4	0	...	0	0	0	0	
4	78.02	20	23	34	5	3	12	3	13	0	...	0	0	0	0	

	X379	X380	X382	X383	X384	X385
0	0	0	0	0	0	0
1	0	0	0	0	0	0
2	0	0	1	0	0	0

```
3    0    0    0    0    0    0
4    0    0    0    0    0    0
```

[5 rows x 365 columns]

```
[47]: df_test.head()
```

```
[47]:   X0  X1  X2  X3  X4  X5  X6  X8  X10  X11  ...  X375  X376  X377  X378  \
0   21  23  17   5   3  24   0  22   0   0  ...    0    0    0    1
1   42   3  19   0   3  28   6  24   0   0  ...    0    0    1    0
2   21  23  34   5   3  27   9   9   0   0  ...    0    0    0    1
3   21  13  34   5   3  27  11  13   0   0  ...    0    0    0    1
4   45  20  34   2   3  12   8  12   0   0  ...    1    0    0    0
```

```
      X379  X380  X382  X383  X384  X385
0         0     0     0     0     0     0
1         0     0     0     0     0     0
2         0     0     0     0     0     0
3         0     0     0     0     0     0
4         0     0     0     0     0     0
```

[5 rows x 371 columns]

0.0.4 Perform dimensionality reduction.

```
[48]: from sklearn.decomposition import PCA
```

```
[49]: pca =PCA(n_components=0.95) #95% of PCA(Principal Component Analysis used to
      ↪reduce the linearity of dimension)
```

```
[50]: pca.fit(df_train)
```

```
[50]: PCA(n_components=0.95)
```

```
[51]: x_train_95 =pca.transform(df_train)
```

```
[52]: print("the train set shape",df_train.shape)
      print("PCA x train data shape",x_train_95.shape)
```

```
the train set shape (4209, 365)
PCA x train data shape (4209, 6)
```


1 PCA for test Dataset 95%

```
[53]: pca.fit(df_test)
      x_test_95=pca.transform(df_test)
```

```
[54]: print("the test set shape",df_test.shape)
      print("PCA x test data shape",x_test_95.shape)
```

```
the test set shape (4209, 371)
PCA x test data shape (4209, 6)
```

we can see the dimensionality linear reduction using PCmethod has been sucessfully done. lets check for pca =98%

```
[55]: pca_98=PCA(n_components =0.98)
      pca_98.fit(df_train)
      x_train_98 =pca_98.transform(df_train)
```

```
[56]: print("the train set shape",df_train.shape)
      print("PCA x train data PCA 95% shape",x_train_95.shape)
      print("PCA x train data PCA 98% shape",x_train_98.shape)
```

```
the train set shape (4209, 365)
PCA x train data PCA 95% shape (4209, 6)
PCA x train data PCA 98% shape (4209, 12)
```

2 PCA for test dtaset 98%

```
[57]: pca_98.fit(df_test)
      x_test_98 =pca_98.transform(df_test)
```

```
[58]: print("the test set shape",df_test.shape)
      print("PCA x test data PCA 95% shape",x_test_95.shape)
      print("PCA x test data PCA 98% shape",x_test_98.shape)
```

```
the test set shape (4209, 371)
PCA x test data PCA 95% shape (4209, 6)
PCA x test data PCA 98% shape (4209, 13)
```

3 Train and test Split method

```
[59]: X =df_train.drop('y',axis=1)
      y= df_train.y
```

```
[60]: X.head()
```

```
[60]:
```

	X0	X1	X2	X3	X4	X5	X6	X8	X10	X12	...	X375	X376	X377	X378	\
0	32	23	17	0	3	24	9	14	0	0	...	0	0	1	0	
1	32	21	19	4	3	28	11	14	0	0	...	1	0	0	0	
2	20	24	34	2	3	27	9	23	0	0	...	0	0	0	0	
3	20	21	34	5	3	27	11	4	0	0	...	0	0	0	0	
4	20	23	34	5	3	12	3	13	0	0	...	0	0	0	0	

	X379	X380	X382	X383	X384	X385
0	0	0	0	0	0	0
1	0	0	0	0	0	0
2	0	0	1	0	0	0
3	0	0	0	0	0	0
4	0	0	0	0	0	0


```
[5 rows x 364 columns]
```

```
[61]: y
```

```
[61]:
```

0	130.81
1	88.53
2	76.26
3	80.62
4	78.02
...	
4204	107.39
4205	108.77
4206	109.22
4207	87.48
4208	110.85

```
Name: y, Length: 4209, dtype: float64
```

we can see that we have split the train dataset for training and testing.

```
[62]: x_train,x_test,y_train,y_test =train_test_split(X,y,random_state=4,test_size=0.
↪3)
```

```
[63]: print("the x train ",x_train.shape)
print("the y train ",y_train.shape)
print("the x test ",x_test.shape)
print("the y test ",y_test.shape)
```

```
the x train (2946, 364)
the y train (2946,)
the x test (1263, 364)
the y test (1263,)
```

let us apply dimensionality reduction on these train and test set

```
[64]: pca_train = PCA(n_components = 0.95)
```

```
[65]: x_train = pca_train.fit_transform(x_train)
x_train
```

```
[65]: array([[ -14.61088357,  -5.44468509,   0.80104129,  18.74714849,
          -5.65930201,   3.21497912],
          [-15.27453709,   3.28027982, -12.94264609,   3.16512933,
          -7.24110937,   4.02606781],
          [  0.78016352,  17.59226159,  -0.93552596,  11.42292115,
          -10.44001634,  -2.02288687],
          ...,
          [ 14.14162758, -16.10180716,  -2.83619905, -12.87851378,
           10.12104015,  -2.31762817],
          [  8.29098768,  22.86626337,  -6.5125823 , -10.48590224,
          -1.19922965,  -0.51539951],
          [ 11.58518722, -12.52254749,  -4.87001745,  14.76108837,
           7.97477676,  -1.51903313]])
```

```
[ ]:
```

```
[66]: pca_test = PCA(n_components = 0.95)
pca_test.fit(x_test)
```

```
[66]: PCA(n_components=0.95)
```

```
[67]: x_test = pca_test.fit_transform(x_test)
x_test
```

```
[67]: array([[ -4.31386696,   0.92430474,   9.39805756,   9.74687656,
          -7.7774354 ,  -2.31105816],
          [-5.03267524,   0.45439079,   5.33359858,   8.75556393,
           3.14899705,  -2.34863649],
          [ 23.33451655,  16.63864391,  -6.68758541, -12.42121238,
          -6.69189448,  -1.66313882],
          ...,
          [ 26.87917614,  14.68405351,   7.26613642,  -1.04794556,
          -8.14630772,   4.17370677],
          [  3.35457377, -11.28939448,  10.24108822,  -5.14665291,
           0.69354433,  -4.30753325],
          [-2.62224304, -14.18635134,  -4.58477466, -13.77830518,
          -3.16576757,   0.67861946]])
```

```
[68]: x_train.shape
```

```
[68]: (2946, 6)
```

```
[69]: x_test.shape
```

```
[69]: (1263, 6)
```

```
[70]: pca_test.explained_variance_
```

```
[70]: array([205.58713329, 111.82314965,  71.21523868,  62.4419506 ,
         49.22480611,   8.46646788])
```

```
[71]: pca_test.explained_variance_ratio_
```

```
[71]: array([0.3864524 , 0.21019956, 0.13386684, 0.11737525, 0.09253033,
         0.01591484])
```

3.1 Perform XGBoost

Predict your test_df values using XGBoost.

```
[72]: from sklearn import svm
      from sklearn import model_selection
      from sklearn.metrics import accuracy_score, r2_score, mean_squared_error, precision_score, confusion_matrix, classification_report
      import xgboost as xgb
```

```
[73]: seed = 7
      num_trees = 30
      kfold = model_selection.KFold(n_splits=10, random_state=7, shuffle=True)
      model = xgb.XGBRegressor(objective='reg:linear', colsample_bytree = 0.3,
      ↪ learning_rate = 0.1,
      ↪ max_depth = 10, alpha = 10, n_estimators=num_trees,
      ↪ random_state=7)
      result = model_selection.cross_val_score(model, X, y, cv=kfold)
```

```
[07:26:51] WARNING: /workspace/src/objective/regression_obj.cu:167: reg:linear
is now deprecated in favor of reg:squarederror.
```

```
[07:26:52] WARNING: /workspace/src/objective/regression_obj.cu:167: reg:linear
is now deprecated in favor of reg:squarederror.
```

```
[07:26:54] WARNING: /workspace/src/objective/regression_obj.cu:167: reg:linear
is now deprecated in favor of reg:squarederror.
```

```
[07:26:55] WARNING: /workspace/src/objective/regression_obj.cu:167: reg:linear
is now deprecated in favor of reg:squarederror.
```

```
[07:26:56] WARNING: /workspace/src/objective/regression_obj.cu:167: reg:linear
is now deprecated in favor of reg:squarederror.
```

```
[07:26:58] WARNING: /workspace/src/objective/regression_obj.cu:167: reg:linear
```

is now deprecated in favor of reg:squarederror.
 [07:26:59] WARNING: /workspace/src/objective/regression_obj.cu:167: reg:linear
 is now deprecated in favor of reg:squarederror.
 [07:27:01] WARNING: /workspace/src/objective/regression_obj.cu:167: reg:linear
 is now deprecated in favor of reg:squarederror.
 [07:27:02] WARNING: /workspace/src/objective/regression_obj.cu:167: reg:linear
 is now deprecated in favor of reg:squarederror.
 [07:27:03] WARNING: /workspace/src/objective/regression_obj.cu:167: reg:linear
 is now deprecated in favor of reg:squarederror.

```
[74]: print(result.mean())
```

0.45283444057913275

lets check the mean squared error and r2_score

```
[75]: model.fit(x_train,y_train)
```

[07:27:04] WARNING: /workspace/src/objective/regression_obj.cu:167: reg:linear
 is now deprecated in favor of reg:squarederror.

```
[75]: XGBRegressor(alpha=10, base_score=0.5, booster=None, colsample_bylevel=1,
  colsample_bynode=1, colsample_bytree=0.3, gamma=0, gpu_id=-1,
  importance_type='gain', interaction_constraints=None,
  learning_rate=0.1, max_delta_step=0, max_depth=10,
  min_child_weight=1, missing=nan, monotone_constraints=None,
  n_estimators=30, n_jobs=0, num_parallel_tree=1,
  objective='reg:linear', random_state=7, reg_alpha=10, reg_lambda=1,
  scale_pos_weight=1, subsample=1, tree_method=None,
  validate_parameters=False, verbosity=None)
```

```
[88]: y_pred =model.predict(x_test)
y_pred
```

```
[88]: array([ 97.59818,  93.51209, 100.91509, ...,  99.22959,  93.56673,
  99.0103 ], dtype=float32)
```

Now let us check the mean squared error and r2_score of the data for df_train.

```
[89]: import math
from math import sqrt
```

```
[90]: print(sqrt(mean_squared_error(y_pred,y_test)))
```

12.069975126714981

```
[91]: print(r2_score(y_pred,y_test))
```

-8.419414020856575

3.2 Conclusion:

we have seen the each module predict approximately 95% testing speed for Mercedes_Benz to scale the safety and reliability of every unique car models of Mercedes_Benz configuration before hits the road. Therefore dimensionality reduction helps to create different small scale models of a very large datasets and then helps to ensemble it for prediction by using XGBoost gradient function which wil further predict the algortithm for robust test procedure which includes faster testing, resulting in lower carbon dioxide emissions without reducing Mercedes-Benz's standards.

[]: